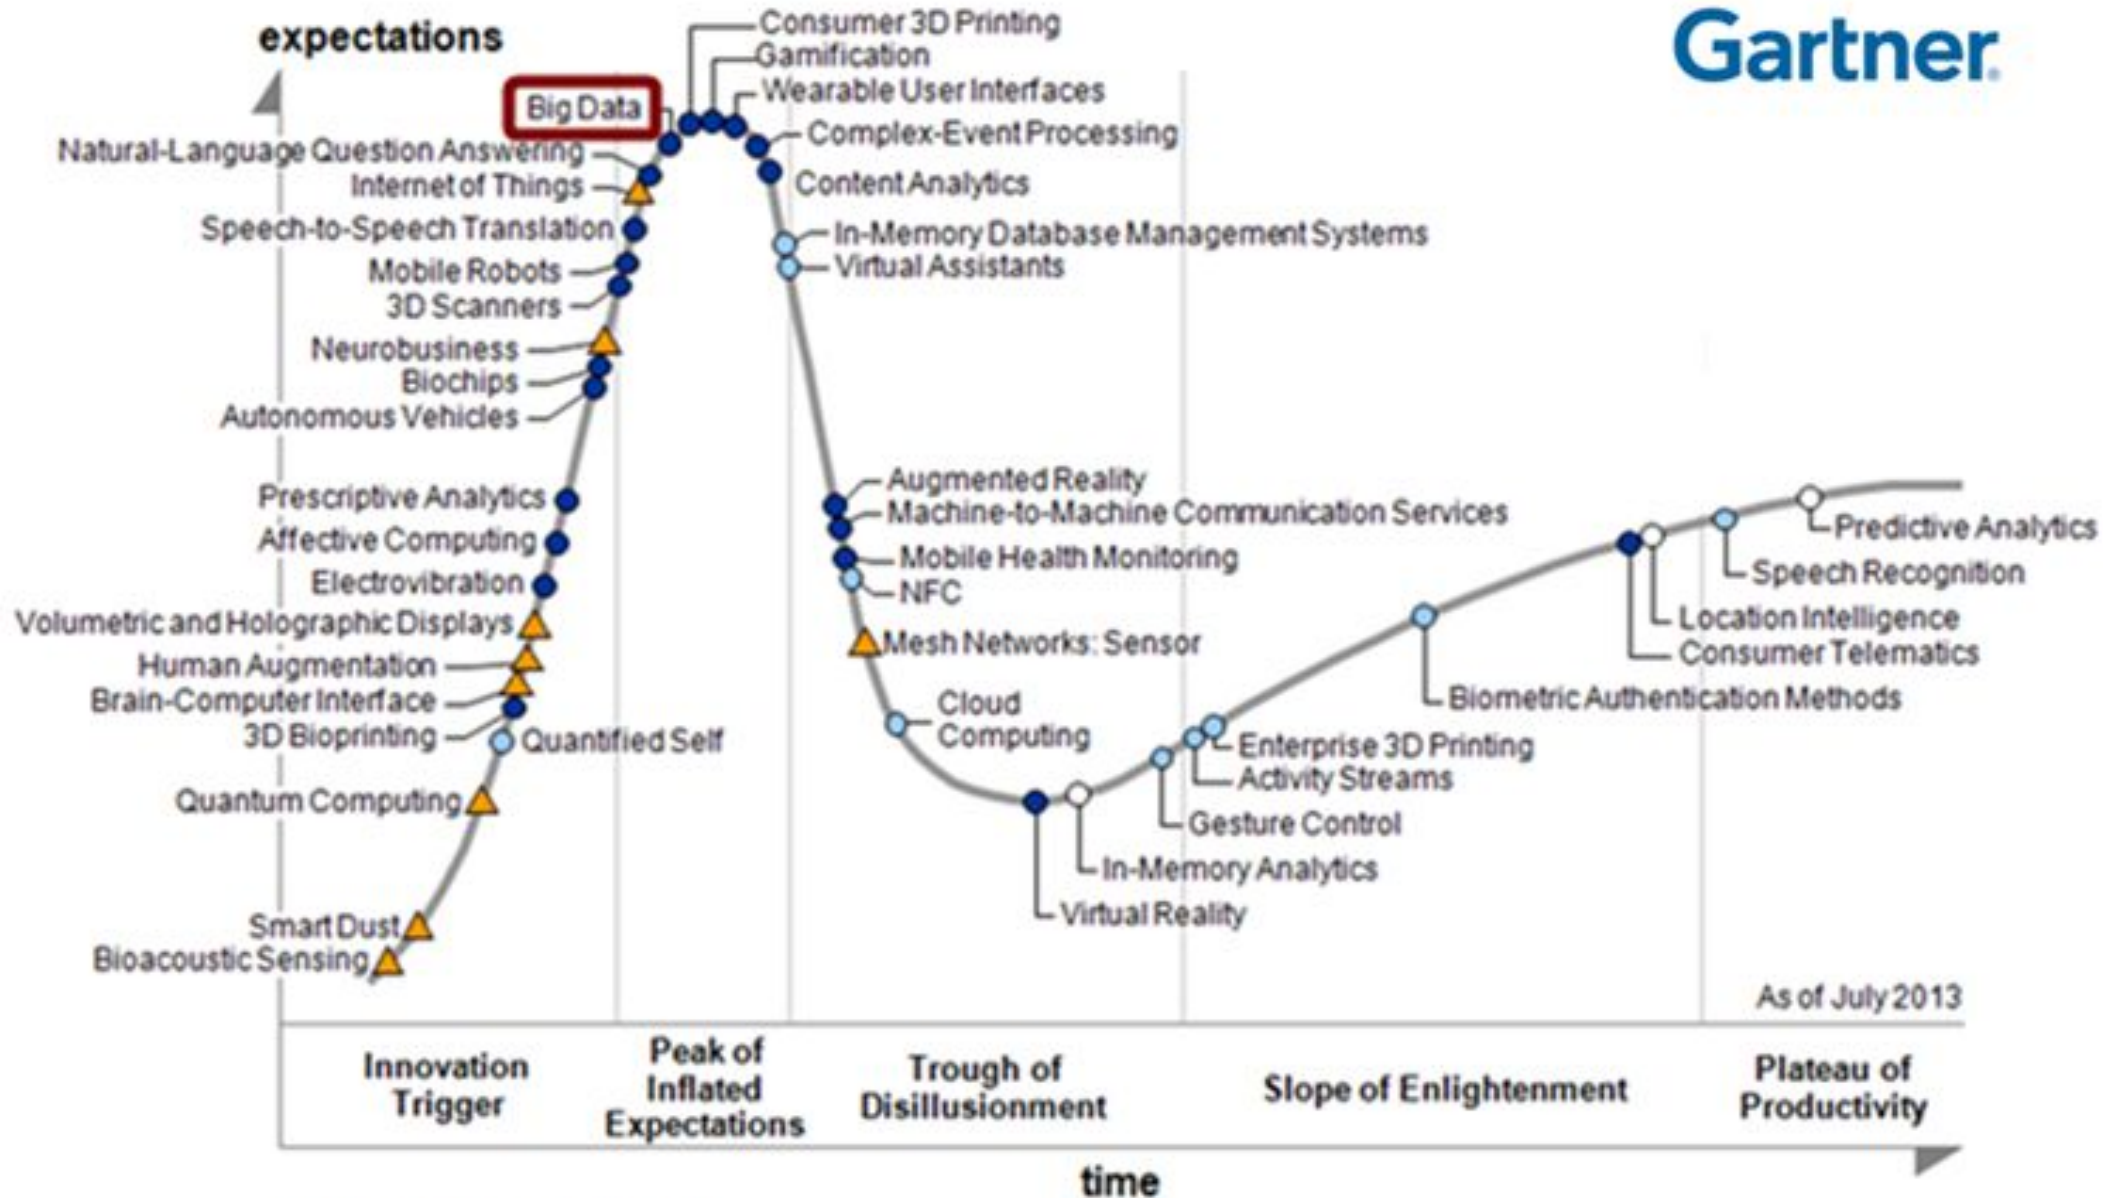


Kimberly Wilber
Sep 4, 2018

“Hype Cycle for Emerging Technologies (2013).”

Gartner



Plateau will be reached in:

○ less than 2 years

● 2 to 5 years

● 5 to 10 years

▲ more than 10 years

○ obsolete
before plateau

Terabytes is not big data, petabytes is

Submitted by **hingo** on Fri, 2011-01-07 22:12

datamining | **MySQL** | **NoSQL**

I often wonder what's behind the increased trend behind Hadoop and other NoSQL technologies. I realize if you're Yahoo that such technology makes sense. I don't get why everyone else wants to use it.



Reading **Stephen O'Grady's self-review of his predictions**

Tw



Ho
up
to
bu
Are

Look guys... **Terabytes is not Big Data!** (And Gigabytes never was big data, even Excel can now take that amount of data :-)



DevOps Borat

@DEVOPS_BORAT

 **Follow**

**Big Data is any thing which is crash
Excel.**

 Reply  Retweet  Favorite  More

RETWEETS

1,977

FAVORITES

423



9:25 AM - 8 Jan 2013

This lecture: Show off two tools for working with large datasets


- **Multiprocessing and parallelism**
 - Chop into pieces
- **Code demo**
- **MapReduce**
 - Add a reduction step
- **When** is it appropriate to use each?

Multiprocessing since ancient times

Septuagint: Translating the Old Testament from Hebrew to Greek in 2nd century BCE



Legend: 72 scholars completed the translation in seventy days.

לְתוֹרָה לָהֶם מִצִּוְתָהּ וְעַצֵּן יִהְיֶה עֲלֵיהֶם יוֹמָם בְּנִסְעָם
מִן הַמִּדְבָּר וְיִהְיֶה בְנִסְעָם
הָאָרֶץ וַיֹּאמֶר מֹשֶׁה קוֹמָה יִהְיֶה וַיַּפְצִּי אֵיבִיךָ וַיָּנֹסוּ
מִשְׁנֵאִיךָ מִפָּנֶיךָ וּבְנִזָּה יֹאמֶר שׁוּבָה יִהְיֶה רַבְבוֹת
אֲלֵפֵי יִשְׂרָאֵל
וְיִהְיֶה הָעָם כְּמֹת אֲנָשִׁים רָעִים בְּאֶזְצִי יִהְיֶה וַיִּשְׁמַע יְהוָה









ΒΑΣΙΛΕΥΣ ΠΕΡΣΩΝ ΚΑΙ
ΕΚ ΗΓΥΣΕΝ ΟΛΗ ΤΗ ΒΑΣΙ
ΛΕΙΑ ΑΥΤΟΥ ΚΑΙ ΕΚ ΜΑΔΙΑ
ΓΡΑΠΤΩΝ ΛΕΓΩΝ ΤΑ
ΔΕ ΛΕΓΕΙ Ο ΒΑΣΙΛΕΥΣ ΠΕΡ
ΣΩΝ ΚΥΡΟΣ ΕΜΕ ΑΝΕΔ

Step 1: Partition




Step 2: Map

Life
sucks!



Step 3: Reduce/Aggregate



When does this work?

Tasks that are “**embarrassingly parallel**” are easy to split up.

This means there should be **no dependencies** between tasks.

- Computing statistics on different slices of the dataset
- Running experiments with different parameters each time

How to use multiprocessing?

There are several good multiprocessing libraries for Python:

- The `multiprocessing` library
- IPython Parallel

Code demo

<https://gist.github.com/gcr/a25d6a0810185a806623b051024673dc>

Large datasets need more than ordinary parallelization!

Adapted to problems at Internet scale

“A typical MapReduce computation processes many terabytes of data on thousands of machines.” (2004)

- Involving **many machines** (network transfer...)
- **Recovering from failures** is important
- **Aggregating results** is tricky
- **Storing data** becomes a problem

“MapReduce: Simplified Data Processing on Large Clusters,” J. Dean and S. Ghemawat, OSDI 2004

“MapReduce”


Multiprocessing is great, but what if you need even more flexibility?

Full-blown **MapReduce** pipelines have several steps:

- **Partition / Split** the data into key/values
- **Map** each datum
- **Sort** the results
- **Reduce** / aggregate the results into the final output


Formal definition

Input:



Example: What are the most frequent words in Shakespeare's plays?

Until very recently, this problem required **multiple man-years** to solve!



Barnardo. Who's there?

Fran. Nay answer me: Stand & unfold
your selfe

Bar. Long live the King

Fran. Barnardo?

Bar. He

Fran. You come most carefully upon
your houre

Bar. 'Tis now strook twelve, get thee
to bed Francisco

Step 1: define the “mapper”

```
def mapper(line):  
    for word in line.split():  
        emit(word, 1)
```

```
mapper("Tis now strook twelve...")
```

```
    {"Tis": 1}
```

```
    {"Now": 1}
```

```
    {"Strook": 1}
```

```
...
```

```
mapper("Gee, what's this ghost doing here?")
```

```
    {"Gee": 1}
```

```
    {"What's": 1}
```

```
...
```

Step 1: define the “mapper”

```
def mapper(line):  
    for word in line.split():  
        emit(word.lower(), 1)
```

```
mapper("Tis now strook twelve...")
```

```
    {"tis": 1}
```

```
    {"now": 1}
```

```
    {"strook": 1}
```

```
    ...
```

```
mapper("Gee, what's this ghost doing here?")
```

```
    {"gee": 1}
```

```
    {"what's": 1}
```

```
    ...
```


Step 2: Sorting

Provided by the framework -- you don't write it!

The sorting step aggregates all results with the same key **together** into a single list.

```
{“tis”: 1}  
{“now”: 1}  
{“strook”: 1}  
{“the”: 1}  
{“twelve”: 1}  
{“romeo”: 1}  
{“the”: 1}
```

...

```
{“tis”: [1,1,1,1,1]}  
{“now”: [1,1,1]}  
{“strook”: [1,1]}  
{“the”: [1,1,1,1,1,1,1,1,1,1]}  
{“twelve”: [1,1]}  
{“romeo”: [1,1,1,1,1,1,1]}  
{“juliet”: [1,1,1,1,1,1,1]}
```

...

Step 3: define the “reducer”

Aggregates all the results together.

```
def reducer(word, counts):  
    emit(word, sum(counts))
```

```
reducer("tis", [1,1,1,1,1])  
    {"tis": 5}  
reducer("the", [1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,  
    {"the": 23590}  
reducer("strook", [1,1])  
    {"strook": 2}  
...
```

Step 3: define the “reducer”

This is incorrect. Can you see why?


```
def reducer(word, counts):  
    emit(word, len(counts))
```

Step 3: define the “reducer”

This is incorrect. Can you see why?

```
def reducer(word, counts):  
    emit(word, len(counts))
```

Reducers run in a **tree**, which means reduce output output may be sent **through other reducers!**




```
reducer("the", [23590, 10201, 330])  
    {"the": 3} # bad!
```

Example 2: An iteration of PageRank

How popular is a certain webpage?

Determined by the number of pages that link to it, and the quality of those pages



We want to compute a score for each page concurrently


```
def mapper(url):  
    webpage = download(url)  
    for outgoing_link in webpage.links:  
        emit(outgoing_link, url)
```

In 1: "cnn.com"

Out 1: → {"wikipedia.org/...": "cnn.com",
 "safeway.com/...": "cnn.com",
 "olivegarden.com/...": "cnn.com"}

In 2: "wikipedia.org"

Out 2: → {"safeway.com/...": "wikipedia.org",
 "starbucks.com/...": "wikipedia.org",
 "cornell.edu/...": "wikipedia.org"}

```
def reducer(dest, sources):  
    weight = 0  
    for source in sources:  
        weight += popularity[source]  
    emit(dest, weight)
```

```
In 1:  {"starbucks.com":  
        ["cnn.com/Starbucks-out-of-coffee.htm",  
         "wikipedia.org/2014-Coffee-Crisis.htm",  
         "cornell.edu/impact-of-bean-growing.htm"]}
```

```
Out 1:  {"starbucks.com": 10 + 8 + 4 }
```

(This is not quite PageRank, but it's similar)

What does a **real** implementation do for you?

- **Machine failures:** With so many machines, $P(\text{no failures})$ is quite small
- **Stragglers at the end of the job:** Performance may depend on the slowest machine. Some implementations resend tasks, etc...
- **Data access and transfer issues:** How can you get TBs of data onto your compute machines?

Do we need MapReduce?

(...Maybe.)

Word frequency as a shell script

```
find -iname '*.txt' \
| xargs cat           \
| tr ' ' '\n'        \
| sort               \
| uniq -c            \
| sort -n            \
| tail
```


Word frequency as a shell script

find -iname '*.txt'	\	Partition step
xargs cat	\	Map step
tr ' ' '\n'	\	
sort	\	Sort step
uniq -c	\	
sort -n	\	Reduce step
tail		

Took a few minutes to write; runs in seconds on a simple laptop.

Why fight for weeks with mapreduce tools?

Adam Drake

Command-line tools can be 235x faster than your Hadoop cluster

Sat 25 January 2014 by [Adam Drake](#)

Introduction

As I was browsing the web and catching up on some sites I visit periodically, I found a cool article from [Tom Hayden](#) about using [Amazon Elastic Map Reduce](#) (EMR) and [mrjob](#) in order to compute some statistics on win/loss ratios for chess games he downloaded from the [millionbase archive](#), and generally have fun with EMR. Since the data volume was only about 1.75GB containing around 2 million chess games, I was skeptical of using Hadoop for the task, but I can understand his goal of learning and having fun with mrjob and EMR

Receive
news

email@example.com

SUBSCRIBE

SITE

[Archives](#)

[Atom feed
\(posts\)](#)

[Atom feed](#)

Jure Leskovec

Organizer of “Mining Large Datasets” MOOC



Jure said every grad student in his lab has one of these machines, and that almost every data set of interest fits in RAM. Contemplate that for a moment.

C.J. Lin

Principal Author of LibSVM

“In my recent visit to a large company, their people did say that most analytics works are still done on one machine.”

Frank McSherry

Microsoft, coauthor of Naiad

“Lots of people struggle with the complexities of getting big data systems up and running, when they possibly shouldn’t be using the systems in the first place. The data sets above are certainly not small (billions of edges), but still run just fine on a laptop. Much faster than the distributed systems, at least.”

Efficient algorithms can be faster

Label propagation to fixed-point (graph connectivity)

System	cores	twitter_rv	uk_2007_05
Spark	128	1784s	8000s+
Giraph	128	200s	8000s+
GraphLab	128	242s	714s
GraphX	128	251s	800s
Single thread	1	153s	417s

Efficient algorithms can be faster

Label propagation to fixed-point (graph connectivity)

System	cores	twitter_rv	uk_2007_05
Single thread (simple)	1	153s	417s
Single thread (smarter)	1	15s	30s

Should we “just split the work,” or do we need all of mapreduce?

Use the right tool for the job. Though the formal model is more general, it comes at a heavy cost and is harder to set up.

Most of the data you interact with will not be big enough to justify the work.

My rule of thumb: Only take the effort if it saves more time than you spend

Sometimes you can also use **better algorithms** to simplify your life. MapReduce feels like brute force.

GraphChi: How a Mac Mini Outperformed a 1,636 Node Hadoop Cluster



Christian Prokopp, Data Scientist, Rangespan

10/29/2013

[Comment](#)


8 comments



[Login](#)

Last year, [GraphChi](#), a spin-off of GraphLab, a distributed graph-based, high-performance computation framework, did something remarkable.

GraphChi outperformed a 1,636 node Hadoop cluster processing a Twitter graph (dataset from 2010) with 1.5 billion edges -- using a single Mac Mini. The task was triangle counting and the Hadoop cluster required over seven days while GraphChi on the Mac Mini did it in one hour! The



When does MapReduce shine?

When you **cannot move data around** – the dataset is too big to copy to your local machine

When the output of the map stage is **much smaller than its input**, or when the reduce phase is non-existent

When **fault-tolerance** is important

When **data storage and transfer** is important